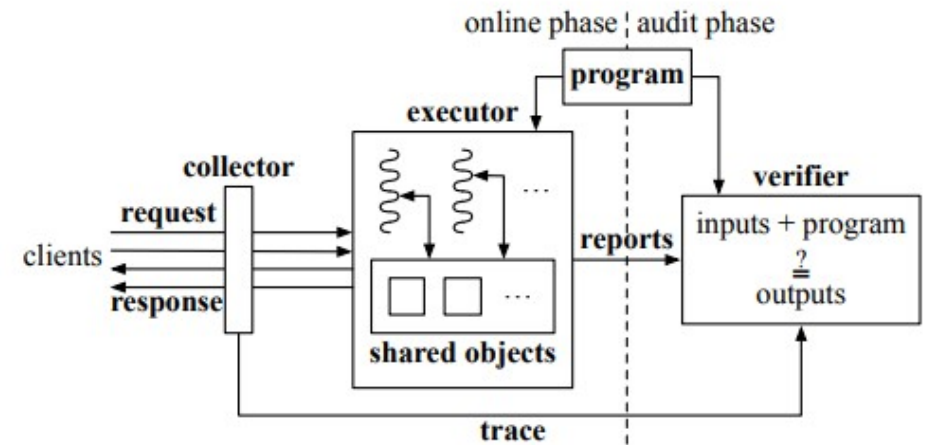


# The Efficient Server Audit Problem, Deduplicated Re-execution, and the Web

Cheng Tan, Lingfan Yu, Joshua B. Leners, and Michael  
Walfish

# The Efficient Server Audit Problem

- Clients issue requests(inputs) to the executor and receive responses(outputs)
- Collector captures trace(accurately)
- Executor maintains reports(untrusted)
- Verifier is responsible for audit process
- Verifier is weaker than the executor
- Executor is permitted to handle multiple requests at the same time
- Shared objects: DB



# The Efficient Server Audit Problem

- Design the verifier and the reports to meet these properties
  - Completeness
    - Verifier must accept the given trace if the executor executed the given program
  - Soundness
    - Verifier must reject if the executor misbehaved during the time period of the trace
  - Efficiency
    - Verifier must require only a small fraction of the computational resources

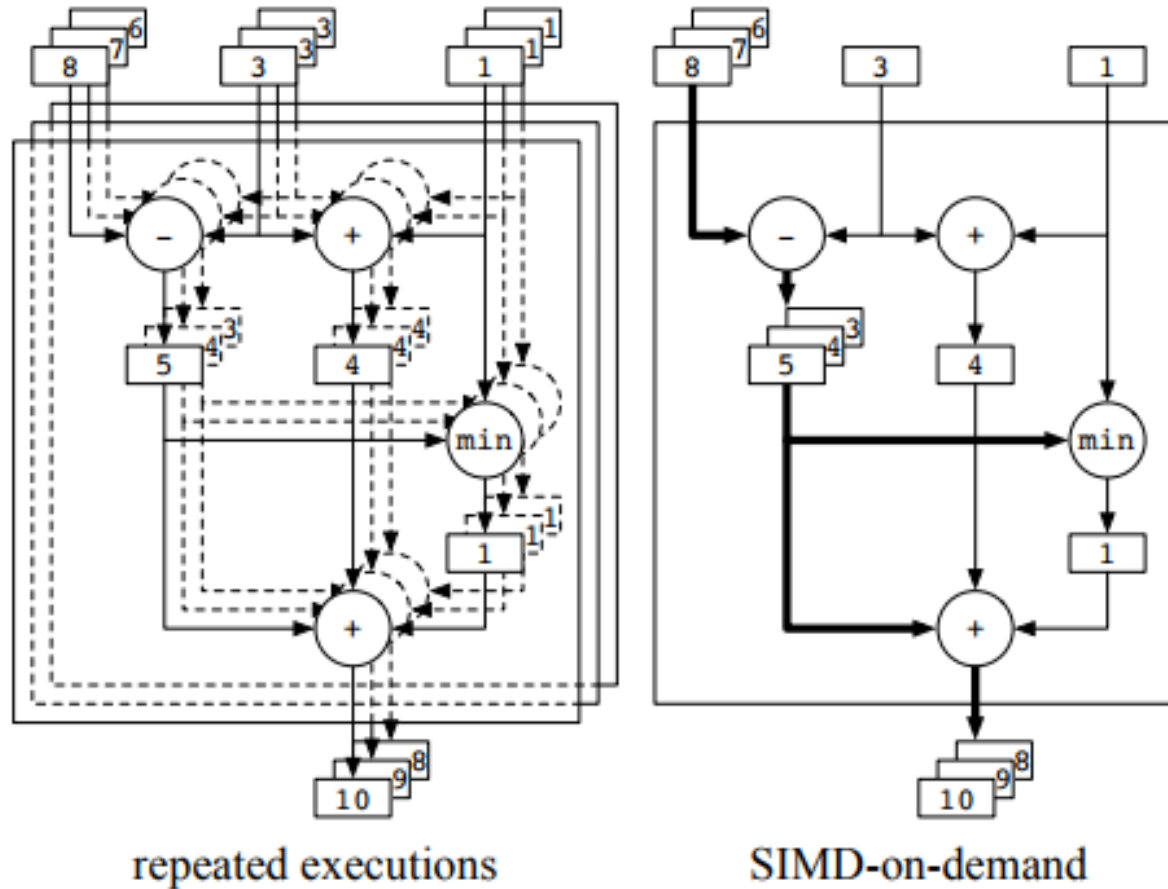
# A Solution:SSCO

- Control flow groupings
  - An opaque tag that purportedly identifies the control flow of the execution
  - Requests that induce the same control flow are supposed to receive the same tag
- Operation logs
  - For each shared object, the executor maintains an ordered log of all operations(across all requests)
- Operation counts
  - For each request execution, the executor records the total number of object operations that is issued.

# SIMD-on-demand execution

- For each control flow group, verifier conducts a single “superposed” execution that logically executes all requests in that group together
- Instructions whose operands are different across the separate logical executions are performed separately
- Instruction executes only once if operands are same

# SIMD-on-demand execution




# Simulate and check

---

Input Trace $Tr$	Input Reports $R$	Global $OpMap: (requestID, opnum) \rightarrow (i, seqnum)$
Components of the reports $R$ :		
$C: CtlFlowTag \rightarrow Set(requestIDs)$ // purported groups; §3.1		24: <b>procedure</b> REEXEC()
$OL_i: \mathbb{N}^+ \rightarrow (requestID, opnum, optype, opcontents)$ // purported op logs; §3.3		25:   Re-execute $Tr$ in groups according to $C$ :
$M: requestID \rightarrow \mathbb{N}$ // purported op counts; §3.3		26:
		27:   (1) Initialize a group as follows:
		28:       Read in inputs for all requests in the group
		29:       Allocate program structures for each request in the group
		30: $opnum \leftarrow 1$ // $opnum$ is a per-group running counter
		31:
		32:   (2) During SIMD-on-demand execution (§3.1):
		33:
		34:       if execution within the group diverges: <b>return</b> REJECT
		35:
		36:       When the group makes a state operation:
		37: $optype \leftarrow$ the type of state operation
		38:        for all $rid$ in the group:
		39: $i, oc \leftarrow$ state op parameters from execution
		40: $s \leftarrow$ CheckOp( $rid, opnum, i, optype, oc$ ) // line 7
		41:         if $optype = \text{RegisterRead}$ :
		42:         state op result $\leftarrow$ SimOp( $i, s, optype, oc$ ) // line 16
		43: $opnum \leftarrow opnum + 1$
		44:
		45:   (3) When a request $rid$ finishes:
		46:       if $opnum < M(rid)$ : <b>return</b> REJECT
		47:
		48:   (4) Write out the produced outputs
		49:
		50:   if the produced outputs from (4) are exactly the responses in $Tr$ :
		51: <b>return</b> ACCEPT
		52: <b>return</b> REJECT
1: <b>procedure</b> SSCO_AUDIT()		
2:   // Partially validate reports (§3.5) and construct $OpMap$		
3:   ProcessOpReports()   // defined in Figure 5		
4:		
5: <b>return</b> ReExec()   // line 24		
6:		
7: <b>procedure</b> CHECKOP( $rid, opnum, i, optype, opcontents$ )		
8:   if ( $rid, opnum$ ) not in $OpMap$ : REJECT		
9:		
10: $\hat{i}, s \leftarrow OpMap[rid, opnum]$		
11: $\hat{o}t, \hat{o}c \leftarrow (OL_i[s].optype, OL_i[s].opcontents)$		
12:   if $i \neq \hat{i}$ or $optype \neq \hat{o}t$ or $opcontents \neq \hat{o}c$ :		
13:     REJECT		
14: <b>return</b> $s$		
15:		
16: <b>procedure</b> SIMOP( $i, s, optype, opcontents$ )		
17: $ret \leftarrow \perp$		
18: $writeop \leftarrow$ walk backward in $OL_i$ from $s$ ; stop when		
19: $optype = \text{RegisterWrite}$		
20:   if $writeop$ doesn't exist:		
21:     REJECT		
22: $ret = writeop.opcontents$		
23: <b>return</b> $ret$		

---

# Simulate and check is not enough





# Consistent ordering verification

---

```
1: procedure CREATETIMEPRECEDENCEGRAPH()
2:   // "Latest" requests; "parent(s)" of any new request
3:   Frontier  $\leftarrow \{\}$ 
4:    $G_{Tr}.Nodes \leftarrow \{\}$ ,  $G_{Tr}.Edges \leftarrow \{\}$ 
5:
6:   for each input and output event in Tr, in time order:
7:     if the event is REQUEST(rid):
8:        $G_{Tr}.Nodes \mathrel{+}= rid$ 
9:       for each r in Frontier:
10:         $G_{Tr}.Edges \mathrel{+}= \langle r, rid \rangle$ 
11:     if the event is RESPONSE(rid):
12:       // rid enters Frontier, evicting its parents
13:        $Frontier \mathrel{-}= \{ r \mid \langle r, rid \rangle \in G_{Tr}.Edges \}$ 
14:        $Frontier \mathrel{+}= rid$ 
15:   return  $G_{Tr}$ 
```

---


---

```
1: Global Trace Tr, Reports R, Graph G, OpMap OpMap
2: procedure PROCESSOPREPORTS()
3:
4:    $G_{Tr} \leftarrow \text{CreateTimePrecedenceGraph}()$  // defined in Figure 6
5:   SplitNodes( $G_{Tr}$ )
6:   AddProgramEdges()
7:
8:   CheckLogs() // also builds the OpMap
9:   AddStateEdges()
10:
11:   if CycleDetect(G): // standard algorithm; see [31, Ch. 22]
12:     REJECT
13:
```

# OROCHI

- Orochi targets apps based on PHP and SQL(LAMP)
- Server and verifier: modified PHP runtimes
- Built atop HipHop VM
- 20K lines of C++, PHP, Bash, Python
- Applications
  - MediaWiki, phpBB and HotCRP

# Orochi's verifier is efficient



# The price of verifiability is tolerable

App	audit speedup	server CPU overhead	avg request	reports (per request)			DB overhead	
				baseline	OROCHI	OROCHI ovhd	temp	permanent
MediaWiki	10.9×	4.7%	7.1KB	0.8KB	1.7KB	11.4%	1.0×	1×
phpBB	5.6×	8.6%	5.7KB	0.1KB	0.3KB	2.7%	1.7×	1×
HotCRP	6.2×	5.9%	3.2KB	0.0KB	0.4KB	10.9%	1.5×	1×

