

Precise and Scalable Detection of Double-Fetch Bugs in OS Kernels

Meng Xu, Chenxiong Qian, Kangjie
Lu, Michael Backes, Taesoo Kim

Background

- Virtual memory is divided into userspace and kernel-space regions
- Userspace memory can be accessed from all threads running in that address space as well as from kernel
- Kernel almost never directly dereferences an address supplied by user processes
- Kernel duplicates the data into kernel memory with transfer functions. (`copy_from_user`, `get_user`)
- (`_user` mark) are placed to ensure userspace memory can be accessed only through transfer functions.

Prior works

- False alerts and missing bugs
- Manually defined patterns can not cover all possible multi-reads
- No attempts to distinguish double-fetch bugs from multi-reads
- No systematic work

Double-fetch bugs

- Multi-read: there are at least two reads from userspace memory
- Overlapped-fetch: The two fetches must cover an overlapped memory region in the userspace
- A relation must exist based on the overlapped regions between the two fetches (control and data dependence)
- Bugs: Cannot prove that the relation established still holds after the second fetch

Examples: Control dependence

```
1 void tls_setsockopt_simplified(char __user *arg) {
2     struct tls_crypto_info header, *full = /* allocated before */;
3
4     // first fetch
5     if (copy_from_user(&header, arg, sizeof(struct tls_crypto_info)))
6         return -EFAULT;
7
8     // protocol check
9     if (header.version != TLS_1_2_VERSION)
10        return -ENOTSUPP;
11
12    // second fetch
13    if (copy_from_user(full, arg,
14        sizeof(struct tls12_crypto_info_aes_gcm_128)))
15        return -EFAULT;
16
17    // BUG: full->version might not be TLS_1_2_VERSION
18    do_something_with(full);
19 }
```

Examples: Data dependence

```
1 void mptctl_simplified(unsigned long arg) {
2     mpt_ioctl_header khdr, __user *uhdr = (void __user *) arg;
3     MPT_ADAPTER *iocp = NULL;
4
5     // first fetch
6     if (copy_from_user(&khdr, uhdr, sizeof(khdr)))
7         return -EFAULT;
8
9     // dependency lookup
10    if (mpt_verify_adapter(khdr.iocnum, &iocp) < 0 || iocp == NULL)
11        return -EFAULT;
12
13    // dependency usage
14    mutex_lock(&iocp->iocctl_cmds.mutex);
15    struct mpt_fw_xfer kfwdl, __user *ufwdl = (void __user *) arg;
16
17    // second fetch
18    if (copy_from_user(&kfwdl, ufwdl, sizeof(struct mpt_fw_xfer)))
19        return -EFAULT;
20
21    // BUG: kfwdl.iocnum might not equal to khdr.iocnum
22    mptctl_do_fw_download(kfwdl.iocnum, .....);
23    mutex_unlock(&iocp->iocctl_cmds.mutex);
24 }
```

Examples: Both

```
1 static int perf_copy_attr_simplified
2 (struct perf_event_attr __user *uattr,
3  struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19
20    .....
21 }
22 // Example: if attr->size is used later
23 // BUG: attr->size can be very large
24 memcpy(buf, attr, attr->size);
```

Formal terms

- (A, S) to denote a fetch. $A \rightarrow$ start address, $S \rightarrow$ size of the memory
- Two fetches: (A_0, S_0) & (A_1, S_1) ; Overlapped memory (A_0, S_0) & (A_1, S_1) ;
 - $A_0 \leq A_1 < A_0 + S_0$ or $A_1 \leq A_0 < A_1 + S_1$
- Control dependence: V' must satisfy V
- Data dependence: $V' == V$

Overview

Algorithm 1: High-level procedure for *double-fetch bug* detection

In : *Kernel* - The kernel to be checked
Out : *Bugs* - The set of *double-fetch bugs* found

```
1 Bugs  $\leftarrow \emptyset$ 
2 Setf  $\leftarrow$  Collect_Fetches(Kernel);
3 for F  $\in$  Setf do
4   Setmr  $\leftarrow$  Collect_Multi_Reads(F)
5   for  $\langle F_0, F_1, F_n \rangle \in$  Setmr do
6     Paths  $\leftarrow$  Construct_Execution_Paths(F0, F1, Fn)
7     for P  $\in$  Paths do
8       if Symbolic_Checking(P, F0, F1) == UNSAFE then
9         Bugs.add( $\langle F_0, F_1 \rangle$ )
10      end
11    end
12  end
13 end
```

Finding multi-reads

- Identify all fetches in the kernel
- Construct a complete, inter-procedural CFG for the whole kernel
- Perform pairwise reachability tests for each pair of fetches


From multi-reads to double-fetch

```

1 static int perf_copy_attr_simplified
2 (struct perf_event_attr __user *uattr,
3  struct perf_event_attr *attr) {
4
5     u32 size;
6
7     // first fetch
8     if (get_user(size, &uattr->size))
9         return -EFAULT;
10
11    // sanity checks
12    if (size > PAGE_SIZE ||
13        size < PERF_ATTR_SIZE_VER0)
14        return -EINVAL;
15
16    // second fetch
17    if (copy_from_user(attr, uattr, size))
18        return -EFAULT;
19    .....
21 }
22 // Example: if attr->size is used later
23 // BUG: attr->size can be very large
24 memcpy(buf, attr, attr->size);

```

(a) C source code



(b) Memory access patterns

```

1 // init root SR
2 $0 = $PARM(0), @0 = $UMEM(0) // uattr
3 $1 = $PARM(1), @1 = $KMEM(0) // attr
4 ---
5 // first fetch
6 fetch(F1) is {A = $0 + 4, S = 4}
7 $2 = @0(4, 7, U0), @2 = nil // size
8 ---
9 // sanity checks
10 assert $2 <= PAGE_SIZE
11 assert $2 >= PERF_ATTR_SIZE_VER0
12 ---
13 // second fetch
14 fetch(F2) is {A = $0, S = $2}
15 @1(0, $2 - 1, K) = @0(0, $2, U1)
16 ---
17 // check fetch overlap
18 assert F2.A <= F1.A < F2.A + F2.S
19     OR F1.A <= F2.A < F1.A + F1.S
20 // --> satisfiable with @0(4, 7, U)
21
22 // check double-fetch bug
23 prove @0(4, 7, U0) == @0(4, 7, U1)
24 // --> fail, no constraints on @0(4, 7, U1)

```

(c) Symbolic representation and checking

Implementation

- Compile source code to LLVM IR
 - Extract build log to collect the compilation flags
 - Use Clang to compile
 - llvm-link to merge

Findings

- 1104 multi-reads
- Confirming previous reported bugs
 - one is miss, due to the IBM S/390 architecture
- New bugs
 - Nine have been fixed
 - four are acknowledged
 - Nine are pending
 - Two won't be fix (Doesn't harm)