

K-Miner: Uncovering Memory Corruption in Linux

David Gens, Simon Schmitt, Lucas
Dava, Ahmad-Reza Sadeghi

Motivation

- static analysis faces severe scalability challenges
- All analysis frameworks for kernel code are limited to intra-procedural analysis
- Linux comprises over 24 million lines of code

Contributions


- Enable global static analysis for kernel code
 - K-Miner(inter-procedural static analysis)
- prototype framework implementation
 - on top of LLVM
- extensive evaluation
 - apply it to all system calls across many different Linux versions

Background


- Data-Flow analysis

```
void *f() {  
1 void *a = alloc();  
2 if (cond(a)) {  
3   free(a);  
4   return NULL;  
5 }  
6 return a;  
}  
  
int main(void) {  
7 void *b = f();  
8 printf("%p\n", b);  
9 return 1;  
}
```


a) Program Code



b) Inter-procedural Control-Flow Graph



c) Pointer Assignment Graph



d) Value-Flow Graph

Figure 1: Data-flow analyses utilize graphs to reason about program behavior at compile time.

Background

- Memory-corruption vulnerabilities
 - integer overflows(IO)
 - use-after-free(UAF)
 - dangling pointers(DP)
 - double free(DF)
 - buffer overflow(BO)
 - missing pointer checks(MPC)

K-Miner

- Assumptions
 - attacker has control user-space
 - OS is isolated from user processes
 - Adversary cannot insert malicious code into the kernel
- Goal
 - Systematically scan the system call interface for these vulnerabilities

Overview

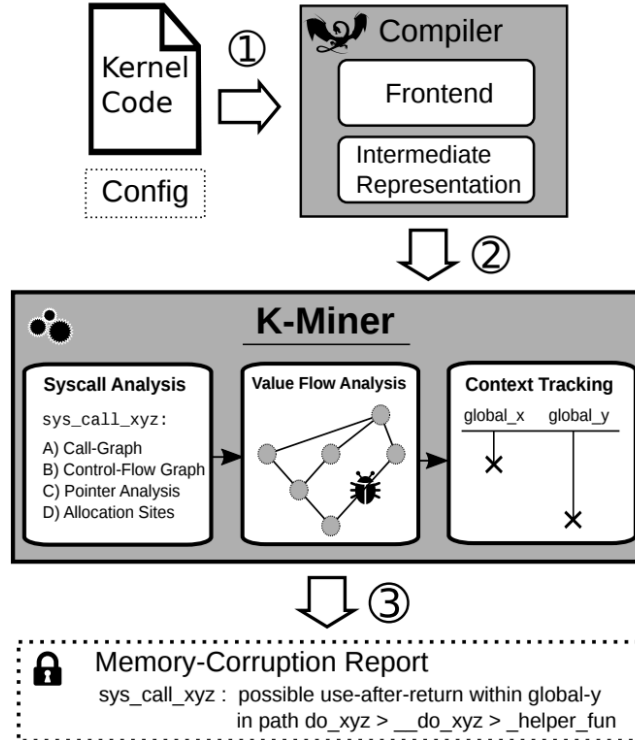


Figure 2: Overview of the different components of K-Miner.

Implementation

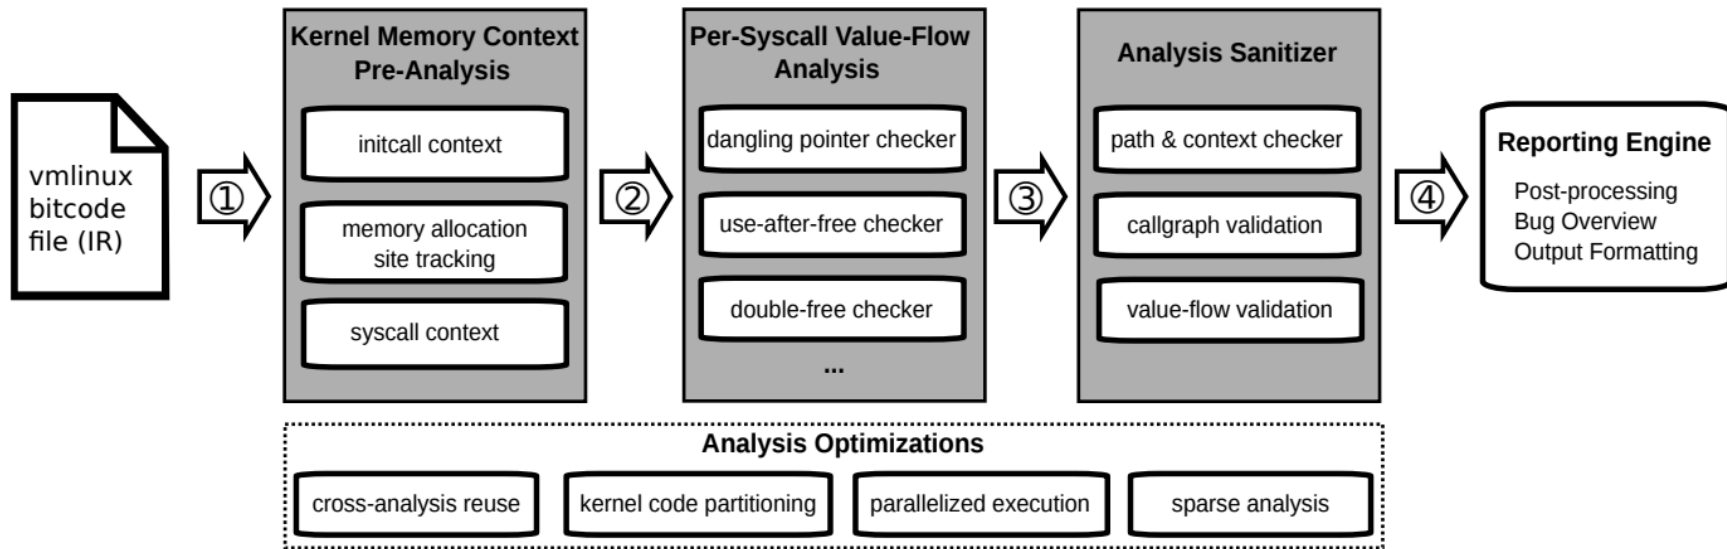


Figure 3: Overview of the K-Miner implementation: we conduct complex data-flow analysis of the Linux kernel in stages, re-using intermediate results.

Evaluation

Version	MLOC	Bitcode	Avg. Run Time	Magnitude of Analysis			Report Results		
				#Functions	#Variables	#Pointers	DP	UAF	DF
3.19	15.5	280M	796.69s	99K	433K	>5M	7 (40)	3 (131)	1 (13)
4.2	16.3	298M	1435.62s	104K	466K	>6M	11 (46)	2 (106)	0 (19)
4.6	17.1	298M	1502.54s	105K	468K	>6M	3 (50)	2 (104)	0 (31)
4.10	22.1	353M	1312.41s	121K	535K	>7M	1 (30)	2 (105)	0 (22)
4.12	24.1	364M	2164.96s	126K	558K	>7.4M	1 (24)	0 (27)	1 (24)

Evaluation




Figure 5: Wall-clock time per analysis phase for system calls requiring more than 30 Minutes within K-Miner.

Evaluation




Figure 6: Maximum memory requirements of K-Miner for system calls requiring more than 11G of RAM.